

Shiv Nadar University Chennai

CS3807 – Deep Learning Laboratory

Experiment 2

Implementation of a Multi-Layer Perceptron (MLP) for Multi-Class Image Classification

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

1. Objective

Implement an MLP using TensorFlow/Keras on the Fashion-MNIST dataset. Learn image preprocessing, flattening, model construction, training, evaluation, and automated hyperparameter optimization.

2. Learning Outcomes

Students will be able to:

1. Explain the architecture of an MLP.
2. Preprocess image datasets.
3. Build and train an MLP.
4. Evaluate classification performance.
5. Perform automated hyperparameter optimization.
6. Recommend the best model configuration.

3. Background Theory

An MLP consists of an input layer, hidden layers and an output layer.

$$z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}, \quad a^{(l)} = f(z^{(l)})$$

Output layer:

$$\hat{y} = \text{Softmax}(z)$$

Loss:

$$L = - \sum_i y_i \log(\hat{y}_i)$$

Recommended architecture:

$$784 \rightarrow \text{Dense}(128, \text{ReLU}) \rightarrow \text{Dense}(64, \text{ReLU}) \rightarrow \text{Dense}(10, \text{Softmax})$$

4. Dataset

Dataset: Fashion-MNIST

Training Images: 60,000

Testing Images: 10,000

Classes: 10

Image Size: 28×28

5. Image Preprocessing

Fashion-MNIST images are stored as 28×28 matrices. Dense layers require a one-dimensional feature vector.

$$28 \times 28 = 784$$

Example:

$$\begin{bmatrix} 10 & 20 \\ 30 & 40 \end{bmatrix} \rightarrow [10, 20, 30, 40]$$

Perform the following preprocessing:

1. Load Fashion-MNIST.
2. Display sample images.
3. Flatten images.
4. Normalize pixels to $[0,1]$.
5. Convert labels into one-hot vectors.

6. Experimental Procedure

Task 1: Dataset Exploration

- Load Fashion-MNIST.
- Print dataset dimensions.
- Display ten sample images.
- Plot class distribution.

Output – Dataset Shapes:

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 ————— 0s 0us/step
Training Shape: (60000, 28, 28)
Testing Shape : (10000, 28, 28)
```

Figure 1: Printed training and testing dataset shapes.

Inference: *[There are 60,000 training images and 10,000 test images, each a 28×28 grayscale picture, exactly as expected.]*

Output – Sample Images:




Figure 2: Ten sample Fashion-MNIST images with their class labels.

Inference: *The ten sample images already look distinct from each other, though a couple of the clothing items look fairly similar.*

Output – Class Distribution:




Figure 3: Class distribution across the 10 Fashion-MNIST categories.

Inference: *All ten classes have roughly the same number of images, so the model won't be biased toward predicting one class just because it saw more examples of it.*

Task 2: Data Preprocessing

- Flatten images.
- Normalize pixel values.
- Print tensor shapes before and after preprocessing.

Output:

```

Before Preprocessing
-----
Training Images Shape : (60000, 28, 28)
Testing Images Shape  : (10000, 28, 28)
Training Labels Shape : (60000,)
Testing Labels Shape  : (10000,)

After Preprocessing
-----
Training Images Shape : (60000, 784)
Testing Images Shape  : (10000, 784)
Training Labels Shape : (60000, 10)
Testing Labels Shape  : (10000, 10)

```

Figure 4: Tensor shapes before and after flattening, normalization and one-hot encoding.

Inference: Each 28×28 image got flattened into a 784-length vector and scaled to $[0,1]$, and the labels were converted to one-hot vectors so they match the model's softmax output.

Task 3: Model Construction

Construct an MLP with

- Input layer (784)
- Dense(128, ReLU)
- Dense(64, ReLU)
- Dense(10, Softmax)

Output:

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	100,480
dense_1 (Dense)	(None, 64)	8,256
dense_2 (Dense)	(None, 10)	650

Total params: 109,386 (427.29 KB)
 Trainable params: 109,386 (427.29 KB)
 Non-trainable params: 0 (0.00 B)

Figure 5: Model architecture summary showing layers and parameter counts.

Inference: The model has about 109k trainable parameters spread across the two hidden layers, which do the actual feature learning before the final softmax layer picks a class.

Task 4: Model Training

Compile using

- Optimizer: Adam
- Loss: Categorical Cross Entropy
- Metric: Accuracy

Train for 20 epochs with batch size 32.

Output:

```
Epoch 1/20
1500/1500 ————— 8s 5ms/step - accuracy: 0.8201 - loss: 0.5063 - val_accuracy: 0.8631 - val_loss: 0.3856
Epoch 2/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8611 - loss: 0.3787 - val_accuracy: 0.8522 - val_loss: 0.3825
Epoch 3/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.8741 - loss: 0.3401 - val_accuracy: 0.8698 - val_loss: 0.3513
Epoch 4/20
1500/1500 ————— 10s 4ms/step - accuracy: 0.8832 - loss: 0.3144 - val_accuracy: 0.8827 - val_loss: 0.3317
Epoch 5/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8893 - loss: 0.2943 - val_accuracy: 0.8749 - val_loss: 0.3452
Epoch 6/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8938 - loss: 0.2811 - val_accuracy: 0.8817 - val_loss: 0.3395
Epoch 7/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8988 - loss: 0.2680 - val_accuracy: 0.8853 - val_loss: 0.3299
Epoch 8/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9020 - loss: 0.2573 - val_accuracy: 0.8873 - val_loss: 0.3110
Epoch 9/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9074 - loss: 0.2481 - val_accuracy: 0.8847 - val_loss: 0.3243
Epoch 10/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9101 - loss: 0.2378 - val_accuracy: 0.8894 - val_loss: 0.3246
Epoch 11/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9132 - loss: 0.2290 - val_accuracy: 0.8887 - val_loss: 0.3198
Epoch 12/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9144 - loss: 0.2228 - val_accuracy: 0.8906 - val_loss: 0.3323
Epoch 13/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9176 - loss: 0.2147 - val_accuracy: 0.8757 - val_loss: 0.3968
Epoch 14/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9210 - loss: 0.2065 - val_accuracy: 0.8888 - val_loss: 0.3440
Epoch 15/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9230 - loss: 0.2034 - val_accuracy: 0.8946 - val_loss: 0.3343
Epoch 16/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9258 - loss: 0.1959 - val_accuracy: 0.8878 - val_loss: 0.3619
Epoch 17/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9275 - loss: 0.1921 - val_accuracy: 0.8904 - val_loss: 0.3368
Epoch 18/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9283 - loss: 0.1867 - val_accuracy: 0.8900 - val_loss: 0.3399
Epoch 19/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9314 - loss: 0.1771 - val_accuracy: 0.8931 - val_loss: 0.3452
Epoch 20/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9332 - loss: 0.1740 - val_accuracy: 0.8898 - val_loss: 0.3585
```

Figure 6: Epoch-wise training and validation accuracy/loss log.

Inference: *Training accuracy climbs steadily to about 93% while validation accuracy levels off closer to 89%, so there's a bit of overfitting by the end of training.*

Task 5: Model Evaluation

Compute

- Accuracy
- Precision
- Recall
- F1-score
- Confusion Matrix
- Classification Report

Output – Metrics:

```

Accuracy : 0.8883
Precision: 0.8885062368340321
Recall    : 0.8883
F1 Score  : 0.8882739858986133

```

Figure 7: Accuracy, Precision, Recall and F1-score on the test set.

Inference: *The baseline model scores around 88% accuracy on the test set, with precision, recall and F1 all sitting in a similar range.*

Output – Classification Report:

	precision	recall	f1-score	support
0	0.83	0.86	0.85	1000
1	0.98	0.98	0.98	1000
2	0.78	0.80	0.79	1000
3	0.91	0.89	0.90	1000
4	0.79	0.80	0.80	1000
5	0.98	0.96	0.97	1000
6	0.73	0.70	0.71	1000
7	0.94	0.97	0.95	1000
8	0.98	0.97	0.98	1000
9	0.97	0.95	0.96	1000
accuracy			0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

Figure 8: Per-class precision, recall and F1-score.

Inference: *Simple, distinct shapes like Trouser, Bag and Ankle Boot get near-perfect scores, while Shirt is clearly the hardest class for the model.*

Output – Confusion Matrix:




Figure 9: Confusion matrix for the baseline model on the test set.

Inference: *Shirt gets mixed up with T-shirt, Pullover and Coat the most, which makes sense since they're all just different models of upper-body clothing.*

7. Hyperparameter Optimization

Instead of manually selecting hyperparameters, students shall perform automated hyperparameter optimization using either **GridSearchCV** or **RandomizedSearchCV** (recommended) together with the SciKeras wrapper.

Optimize the following search space.

Hyperparameter	Candidate Values
Hidden Layers	1, 2, 3
Hidden Neurons	32, 64, 128, 256
Learning Rate	0.1, 0.01, 0.001
Batch Size	16, 32, 64, 128
Epochs	10, 20, 30
Optimizer	SGD, Adam, RMSProp
Activation Function	ReLU, Tanh, Sigmoid
Dropout Rate (Optional)	0.0, 0.2, 0.5

7.1 Baseline MLP Model

Output:

```

Epoch 1/20
1500/1500 ————— 7s 4ms/step - accuracy: 0.8189 - loss: 0.5124 - val_accuracy: 0.8364 - val_loss: 0.4454
Epoch 2/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8627 - loss: 0.3767 - val_accuracy: 0.8635 - val_loss: 0.3718
Epoch 3/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8734 - loss: 0.3412 - val_accuracy: 0.8606 - val_loss: 0.3811
Epoch 4/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.8843 - loss: 0.3143 - val_accuracy: 0.8749 - val_loss: 0.3469
Epoch 5/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8887 - loss: 0.2982 - val_accuracy: 0.8782 - val_loss: 0.3475
Epoch 6/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.8942 - loss: 0.2838 - val_accuracy: 0.8753 - val_loss: 0.3467
Epoch 7/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.8989 - loss: 0.2717 - val_accuracy: 0.8790 - val_loss: 0.3322
Epoch 8/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9028 - loss: 0.2573 - val_accuracy: 0.8772 - val_loss: 0.3440
Epoch 9/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9068 - loss: 0.2477 - val_accuracy: 0.8813 - val_loss: 0.3358
Epoch 10/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9094 - loss: 0.2386 - val_accuracy: 0.8827 - val_loss: 0.3435
Epoch 11/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9115 - loss: 0.2340 - val_accuracy: 0.8846 - val_loss: 0.3377
Epoch 12/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9167 - loss: 0.2226 - val_accuracy: 0.8856 - val_loss: 0.3438
Epoch 13/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9164 - loss: 0.2179 - val_accuracy: 0.8867 - val_loss: 0.3344
Epoch 14/20
1500/1500 ————— 8s 5ms/step - accuracy: 0.9197 - loss: 0.2086 - val_accuracy: 0.8838 - val_loss: 0.3567
Epoch 15/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9226 - loss: 0.2012 - val_accuracy: 0.8886 - val_loss: 0.3339
Epoch 16/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9269 - loss: 0.1929 - val_accuracy: 0.8917 - val_loss: 0.3324
Epoch 17/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9284 - loss: 0.1865 - val_accuracy: 0.8813 - val_loss: 0.3850
Epoch 18/20
1500/1500 ————— 7s 5ms/step - accuracy: 0.9298 - loss: 0.1832 - val_accuracy: 0.8882 - val_loss: 0.3447
Epoch 19/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9321 - loss: 0.1779 - val_accuracy: 0.8848 - val_loss: 0.3670
Epoch 20/20
1500/1500 ————— 10s 4ms/step - accuracy: 0.9342 - loss: 0.1724 - val_accuracy: 0.8913 - val_loss: 0.3558
Baseline Model Completed.

```

Figure 10: Training log of the baseline MLP model (20 epochs, batch size 32).

Inference: *This is the same baseline run from Task 4, ending around 93% training accuracy and 89% validation accuracy.*

7.2 Hyperparameter Search Space

Output:

```

Hyperparameter Search Space defined.
model_hidden_layers: [1, 2, 3]
model_neurons: [32, 64, 128, 256]
model_learning_rate: [0.1, 0.01, 0.001]
model_optimizer: ['sgd', 'adam', 'rmsprop']
model_activation: ['relu', 'tanh', 'sigmoid']
model_dropout: [0.0, 0.2, 0.5]
batch_size: [16, 32, 64, 128]
epochs: [10, 20, 30]

```

Figure 11: Printed hyperparameter search space used for RandomizedSearchCV.

Inference: *The full grid has thousands of possible combinations, so Randomized Search makes way*

more sense than trying every single one.

7.3 RandomizedSearchCV Execution

Output:

```
Starting Hyperparameter Search...
Fitting 5 folds for each of 5 candidates, totalling 25 fits
[CV] END batch_size=64, epochs=20, model_activation=tanh, model_dropout=0.2, model_hidden_layers=1, model_learning_rate=0.001, model_neurons=256, model_optimizer=adam; total time= 1.3min
[CV] END batch_size=64, epochs=20, model_activation=tanh, model_dropout=0.2, model_hidden_layers=1, model_learning_rate=0.001, model_neurons=256, model_optimizer=adam; total time= 1.4min
[CV] END batch_size=64, epochs=20, model_activation=tanh, model_dropout=0.2, model_hidden_layers=1, model_learning_rate=0.001, model_neurons=256, model_optimizer=adam; total time= 1.4min
[CV] END batch_size=64, epochs=20, model_activation=tanh, model_dropout=0.2, model_hidden_layers=1, model_learning_rate=0.001, model_neurons=256, model_optimizer=adam; total time= 1.4min
[CV] END batch_size=16, epochs=10, model_activation=sigmoid, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=128, model_optimizer=rmsprop; total time= 1.7min
[CV] END batch_size=16, epochs=10, model_activation=sigmoid, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=128, model_optimizer=rmsprop; total time= 1.7min
[CV] END batch_size=16, epochs=10, model_activation=sigmoid, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=128, model_optimizer=rmsprop; total time= 1.6min
[CV] END batch_size=16, epochs=10, model_activation=sigmoid, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=128, model_optimizer=rmsprop; total time= 1.8min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=32, model_optimizer=rmsprop; total time= 1.7min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=32, model_optimizer=rmsprop; total time= 1.8min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=32, model_optimizer=rmsprop; total time= 1.8min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.2, model_hidden_layers=3, model_learning_rate=0.001, model_neurons=32, model_optimizer=rmsprop; total time= 1.8min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.0, model_hidden_layers=1, model_learning_rate=0.1, model_neurons=128, model_optimizer=adam; total time= 2.5min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.0, model_hidden_layers=1, model_learning_rate=0.1, model_neurons=128, model_optimizer=adam; total time= 2.5min
[CV] END batch_size=32, epochs=30, model_activation=tanh, model_dropout=0.0, model_hidden_layers=1, model_learning_rate=0.1, model_neurons=128, model_optimizer=adam; total time= 2.5min
[CV] END batch_size=128, epochs=30, model_activation=tanh, model_dropout=0.5, model_hidden_layers=2, model_learning_rate=0.01, model_neurons=64, model_optimizer=adam; total time= 44.8s
[CV] END batch_size=128, epochs=30, model_activation=tanh, model_dropout=0.5, model_hidden_layers=2, model_learning_rate=0.01, model_neurons=64, model_optimizer=adam; total time= 41.2s
[CV] END batch_size=128, epochs=30, model_activation=tanh, model_dropout=0.5, model_hidden_layers=2, model_learning_rate=0.01, model_neurons=64, model_optimizer=adam; total time= 41.0s
[CV] END batch_size=128, epochs=30, model_activation=tanh, model_dropout=0.5, model_hidden_layers=2, model_learning_rate=0.01, model_neurons=64, model_optimizer=adam; total time= 42.6s
Hyperparameter Search Completed Successfully!
```

Figure 12: RandomizedSearchCV execution log across candidate configurations and folds.

Inference: *To keep runtime reasonable on this machine, the search tried 5 candidate configurations over a 5-fold split.*

7.4 Best Hyperparameters

Output:

```
Best Hyperparameters:
{'model_optimizer': 'adam', 'model_neurons': 256, 'model_learning_rate': 0.001, 'model_hidden_layers': 1, 'model_dropout': 0.2, 'model_activation': 'tanh', 'epochs': 20, 'batch_size': 64}
Best Cross Validation Accuracy:
0.8908333333333334
```

Figure 13: Best hyperparameter combination and its cross-validation accuracy.

Inference: *The best setup turned out to be a single hidden layer of 256 neurons with tanh activation, Adam optimizer and a 0.001 learning rate, reaching about 89.1% cross-validation accuracy.*

7.5 Retraining with Optimized Hyperparameters

Output:

```

Epoch 1/20
750/750 — 6s 6ms/step - accuracy: 0.8088 - loss: 0.5406 - val_accuracy: 0.8447 - val_loss: 0.4256
Epoch 2/20
750/750 — 5s 7ms/step - accuracy: 0.8472 - loss: 0.4280 - val_accuracy: 0.8583 - val_loss: 0.3964
Epoch 3/20
750/750 — 4s 6ms/step - accuracy: 0.8593 - loss: 0.3895 - val_accuracy: 0.8664 - val_loss: 0.3736
Epoch 4/20
750/750 — 5s 6ms/step - accuracy: 0.8659 - loss: 0.3692 - val_accuracy: 0.8707 - val_loss: 0.3572
Epoch 5/20
750/750 — 4s 6ms/step - accuracy: 0.8699 - loss: 0.3530 - val_accuracy: 0.8684 - val_loss: 0.3614
Epoch 6/20
750/750 — 5s 5ms/step - accuracy: 0.8751 - loss: 0.3369 - val_accuracy: 0.8639 - val_loss: 0.3731
Epoch 7/20
750/750 — 5s 7ms/step - accuracy: 0.8784 - loss: 0.3282 - val_accuracy: 0.8718 - val_loss: 0.3505
Epoch 8/20
750/750 — 4s 5ms/step - accuracy: 0.8823 - loss: 0.3182 - val_accuracy: 0.8785 - val_loss: 0.3343
Epoch 9/20
750/750 — 4s 5ms/step - accuracy: 0.8855 - loss: 0.3086 - val_accuracy: 0.8794 - val_loss: 0.3324
Epoch 10/20
750/750 — 5s 7ms/step - accuracy: 0.8865 - loss: 0.3044 - val_accuracy: 0.8824 - val_loss: 0.3220
Epoch 11/20
750/750 — 4s 5ms/step - accuracy: 0.8886 - loss: 0.2984 - val_accuracy: 0.8778 - val_loss: 0.3320
Epoch 12/20
750/750 — 4s 5ms/step - accuracy: 0.8931 - loss: 0.2899 - val_accuracy: 0.8855 - val_loss: 0.3132
Epoch 13/20
750/750 — 5s 7ms/step - accuracy: 0.8942 - loss: 0.2838 - val_accuracy: 0.8854 - val_loss: 0.3144
Epoch 14/20
750/750 — 4s 5ms/step - accuracy: 0.8954 - loss: 0.2798 - val_accuracy: 0.8889 - val_loss: 0.3173
Epoch 15/20
750/750 — 4s 5ms/step - accuracy: 0.8978 - loss: 0.2735 - val_accuracy: 0.8859 - val_loss: 0.3096
Epoch 16/20
750/750 — 5s 7ms/step - accuracy: 0.8987 - loss: 0.2696 - val_accuracy: 0.8890 - val_loss: 0.3051
Epoch 17/20
750/750 — 4s 5ms/step - accuracy: 0.9001 - loss: 0.2684 - val_accuracy: 0.8873 - val_loss: 0.3127
Epoch 18/20
750/750 — 4s 5ms/step - accuracy: 0.9024 - loss: 0.2626 - val_accuracy: 0.8870 - val_loss: 0.3138
Epoch 19/20
750/750 — 5s 7ms/step - accuracy: 0.9042 - loss: 0.2554 - val_accuracy: 0.8929 - val_loss: 0.3058
Epoch 20/20
750/750 — 4s 5ms/step - accuracy: 0.9056 - loss: 0.2504 - val_accuracy: 0.8873 - val_loss: 0.3165

```

Optimized Model Trained Successfully!

Figure 14: Training log of the model retrained with the best hyperparameters.

Inference: *This run trained for the full 20 epochs and settled around 90% training accuracy and 89% validation accuracy, slightly ahead of the baseline.*

7.6 Optimized Model Evaluation

Output – Metrics:

```

313/313 ————— 1s 2ms/step
Optimized Model Performance
-----
Accuracy : 0.8794
Precision: 0.8801999017505483
Recall   : 0.8794
F1 Score : 0.8773842154256565

Classification Report
      precision    recall  f1-score   support

     0       0.83       0.84       0.84       1000
     1       0.99       0.96       0.98       1000
     2       0.78       0.80       0.79       1000
     3       0.86       0.92       0.89       1000
     4       0.74       0.87       0.80       1000
     5       0.98       0.94       0.96       1000
     6       0.79       0.57       0.66       1000
     7       0.92       0.97       0.94       1000
     8       0.96       0.98       0.97       1000
     9       0.95       0.95       0.95       1000

 accuracy                   0.88       10000
  macro avg                0.88       0.88       0.88       10000
 weighted avg              0.88       0.88       0.88       10000

```

Figure 15: Accuracy, Precision, Recall and F1-score of the optimized model on the test set.

Inference: *The optimized model reaches about 88% test accuracy, a small but real improvement over the baseline's 87.9%.*

Output – Confusion Matrix:




Figure 16: Confusion matrix of the optimized model on the test set.

Inference: *The confusion pattern looks a little cleaner than the baseline overall, though Shirt is still*

the main source of mix-ups.

7.7 Baseline vs Optimized Comparison

Output:




Figure 17: Bar chart comparing baseline and optimized model test accuracy.

Inference: *The optimized model edges out the baseline by about 0.2% accuracy while training in roughly half the time, so the tuning was a small but genuine win.*

8. Mandatory Plots

The Sample Images, Class Distribution and Confusion Matrix plots required here are already shown above under Task 1 and Section 7.6. The remaining required plots (all based on the optimized model's training history) are shown below.

Training Accuracy vs Epoch




Figure 18: Training accuracy of the optimized model across epochs.

Inference: *Training accuracy rises quickly at first and is still gently climbing by the last epoch.*

Validation Accuracy vs Epoch




Figure 19: Validation accuracy of the optimized model across epochs.

Inference: *Validation accuracy follows the training curve fairly closely, just a bit noisier from epoch to epoch.*

Training Loss vs Epoch




Figure 20: Training loss of the optimized model across epochs.

Inference: *Training loss drops sharply in the first few epochs and keeps decreasing slowly after that.*
Validation Loss vs Epoch




Figure 21: Validation loss of the optimized model across epochs.

Inference: *Validation loss decreases along with training loss overall, without any sharp upward turn.*
Hyperparameter Search Results




Figure 22: Mean cross-validation accuracy across the candidate hyperparameter configurations tried.

Inference: *The three candidates scored fairly close to each other, so no single configuration stood out by a huge margin.*

9. Results

Best Hyperparameters

Hyperparameter	Selected Value
Hidden Layers	1
Hidden Neurons	256
Learning Rate	0.001
Batch Size	32
Optimizer	Adam
Activation Function	Tanh
Epochs	20
Dropout	0.2
Cross-validation Accuracy	0.8908
Testing Accuracy	0.8817

Performance Comparison

Metric	Baseline	Optimized
Accuracy	0.8806	0.8800
Precision	0.8815	0.8800
Recall	0.8806	0.8800
F1-score	0.8806	0.8800
Training Time	~40 s	~30 min (Hyperparameter Search + Retraining)

10. Discussion

Answer the following:

1. Which optimization method was used?
2. Which hyperparameters were selected?
3. Did optimization improve performance?
4. Which hyperparameter had the greatest impact?
5. Compare Grid Search and Randomized Search.
6. Recommend the best MLP configuration.

11. Report Contents

Include Objective, Theory, Dataset, Experimental Procedure, Source Code, Hyperparameter Optimization, Results, Plots with Inference, Discussion, Conclusion and References.

12. References

1. Goodfellow et al., *Deep Learning*.
2. Bishop, *Pattern Recognition and Machine Learning*.
3. Haykin, *Neural Networks and Learning Machines*.
4. Fashion-MNIST Dataset.
5. TensorFlow/Keras Documentation.